

A Model For the Time Complexity of the Time Stepper w/ MPI

Max Ruth

```
$Assumptions = { $\tau \in \text{Integers}$ ,  $\tau > 0$ };  
(* Assumptions about the variables we are working with. In particular,  
we want any Sums to evaluate correctly*)
```

Description of the problem

There are a few parameters that we have to keep track of in this problem. They are:

```
Clear[Nx, Ny, L, M, T]  
Nx; Ny; (* Number of processors in the x and y directions *)  
L; (* Length of the x or y direction of the domain (it's square!) *)  
M; (* Number of grid points the x or y direction *)  
T; (* Total time for which the problem is run *)  
 $\tau$ ; (* Number of time steps taken per MPI Communication *)  
CFL; (* The CFL Condition *)  
h; u; v; (*The things we are solving for: height, x velocity, y velocity *)  
 $\alpha$ ;  $\beta$ ; (* Coefficients in the  $\alpha$ -  
     $\beta$  model of communication in MPI. Assume  $\beta$  has units that work with floats naturally *)  
 $\gamma$ ; (* The time required for one predict-correct time step at a single point *)  
 $\delta$ ; (* The time to calculate the speed at a single point *)
```

Derived from these parameters, we can find a few important things:

```
dx =  $\frac{L}{M}$ ; (* The distance between grid points *)  
dt =  $\frac{\text{CFL}}{\text{Max}[u, v]}$  dx; (* The time step (I think. It's something like this) *)  
mx =  $\frac{M}{Nx}$ ; my =  $\frac{M}{Ny}$ ;  
(* Number of grid points belonging to each cell (ignoring any integer nonsense) *)  
mg =  $4\tau$ ; (* Number of ghost grid cells on one side in one dimension *)  
Nt; (* Number of time steps taken. This isn't know a priori,  
but it is independent of the number of processors (assuming dx is fixed). However,  
it does depend on the CFL condition and dx, so it is in the "derived" category. *)
```

The algorithm that we are performing goes like this:

```

for i = 1 : Nt/τ
  Communicate ghost cells
  Find local time step
  Allreduce for global time step
  Perform the time step operation 2 τ times

```

To analyze the complexity of the algorithm, we need to think about the cost of each of these operations individually

Finding Time Step Costs

Here, we calculate the cost of doing a block of time steps **per processor**.

1. Ghost Cell Communication Cost

The variable we are interested in is:

```
cg; (* Cost of ghost cell communication *)
```

To perform the ghost cell communication, we have to communicate with all four neighbors. The first two communications are to the left and right, and communicate blocks of size $ng \times mx$. The last two communications to the top and bottom have to include the corners as well, so they communicate blocks of size $ng \times (my + ng)$. Plugging this into the α - β model, we have

```
In[27]:= cg = 2 (α + β mg mx) + 2 (α + β mg (mg + my)) // Collect[#, {α, β}, FullSimplify] &
```

```
Out[27]:= 4 α + 8 β τ (M (1/Nx + 1/Ny) + 4 τ)
```

2. Find Local Time Step

The cost of this is simply

```
In[43]:= clts = δ mx my;
```

3. Allreduce for Global Time Step

This is the term that has the worst scaling with the number of processors. Assuming a reasonable algorithm, this should be done via some sort of divide-and-conquer, giving an approximate cost

```
In[45]:= car = α Log[Nx Ny];
```

4. Perform the time-step operation 2τ times

Using the cost to predict-correct for a single node γ , we can find the cost to calculate a layer to be

```
In[37]:= clayer =  $\gamma$  (mx + 2 k) (my + 2 k);
```

where k is the number of ghost cells that are being calculated for the next step. That is, the first step after communicating, we have to calculate the value at $mg - 2$ ghost cells in each direction. After that step, we calculate $mg - 4$, and so on until the final step where we calculate 0 ghost cells. In total, this becomes the sum

```
In[38]:= cts = Sum[clayer, {k, 0, mg - 2, 2}] // Collect[#,  $\tau$ , Simplify] &
```

$$\text{Out[38]= } \frac{2 \left(3 M^2 + 8 N_x N_y - 6 M (N_x + N_y) \right) \gamma \tau}{3 N_x N_y} + 8 \left(-4 + M \left(\frac{1}{N_x} + \frac{1}{N_y} \right) \right) \gamma \tau^2 + \frac{128 \gamma \tau^3}{3}$$

As an example, if we communicate every time step ($\tau = 1$), we have:

```
In[39]:= cts /.  $\tau \rightarrow 1$  // FullSimplify
```

$$\text{Out[39]= } \frac{2 \left(M^2 + 8 N_x N_y + 2 M (N_x + N_y) \right) \gamma}{N_x N_y}$$

Finding Total Cost and Average Cost per Time Step

The total cost of the algorithm comes to multiplying the cost per time step block by the number of blocks, i.e.

```
In[46]:= Ctotal =  $\frac{Nt}{\tau}$  (cg + clts + car + cts) // Collect[#,  $\tau$ , FullSimplify] &
```

$$\text{Out[46]= } \frac{2 Nt \left(6 M (N_x + N_y) \left(2 \beta - \gamma \right) + 3 M^2 \gamma + 8 N_x N_y \gamma \right)}{3 N_x N_y} +$$

$$Nt \left(32 \beta + 8 \left(-4 + M \left(\frac{1}{N_x} + \frac{1}{N_y} \right) \right) \gamma \right) \tau + \frac{128}{3} Nt \gamma \tau^2 + \frac{Nt \left(4 \alpha + \frac{M^2 \delta}{N_x N_y} + \alpha \text{Log}[N_x N_y] \right)}{\tau}$$

The average cost per time step would then simply be

```
In[47]:= cavg =  $\frac{Ctotal}{Nt}$  // Collect[#,  $\tau$ , FullSimplify] &
```

$$\text{Out[47]= } \frac{2 \left(6 M (N_x + N_y) \left(2 \beta - \gamma \right) + 3 M^2 \gamma + 8 N_x N_y \gamma \right)}{3 N_x N_y} +$$

$$\left(32 \beta + 8 \left(-4 + M \left(\frac{1}{N_x} + \frac{1}{N_y} \right) \right) \gamma \right) \tau + \frac{128 \gamma \tau^2}{3} + \frac{4 \alpha + \frac{M^2 \delta}{N_x N_y} + \alpha \text{Log}[N_x N_y]}{\tau}$$

Analysis

What is the best τ ?

It seems that we should probably ignore terms that are constant in τ , as these cannot be tuned away via τ . So,

```
In[48]:= cavganalysis =
  cavg - ((D[cav, \tau] // FullSimplify)) /. \tau -> 0 // Collect[#, \tau, FullSimplify] &
```

$$\text{Out[48]= } \left(32 \beta + 8 \left(-4 + M \left(\frac{1}{N_x} + \frac{1}{N_y} \right) \right) \gamma \right) \tau + \frac{128 \gamma \tau^2}{3} + \frac{4 \alpha + \frac{M^2 \delta}{N_x N_y} + \alpha \text{Log}[N_x N_y]}{\tau}$$

Translating this big-O notation, we have

```
In[50]:= cavgtBigO = O[\tau^2] + O[\frac{1}{\tau}];
```

Because the coefficients of both of these terms are positive, there must be some optimal τ that minimizes cavganalysis. The balance is essentially against latency type costs (i.e. terms with α in them and the time to calculate the time step everywhere in the block) and the extra work required to calculate in the ghost cells (scaling for big τ in the τ^2 term). The actual minimum will depend on all of these parameters however, so it is not so pretty to calculate.

Weak Scaling

For the scaling problems, let's assume that $N_x = N_y$, τ is fixed, and that we only care about the cost of a time step (and not of the total cost). That is,

```
In[52]:= cscaling = cavg /. Ny -> Nx
```

$$\text{Out[52]= } \frac{2 (12 M N_x (2 \beta - \gamma) + 3 M^2 \gamma + 8 N_x^2 \gamma)}{3 N_x^2} + \left(32 \beta + 8 \left(-4 + \frac{2 M}{N_x} \right) \gamma \right) \tau + \frac{128 \gamma \tau^2}{3} + \frac{4 \alpha + \frac{M^2 \delta}{N_x^2} + \alpha \text{Log}[N_x^2]}{\tau}$$

For weak scaling, we go one step farther, and assume that $M = M_0 N_x$:

```
In[54]:= cweak = cscaling /. M -> M0 Nx // FullSimplify
```

$$\text{Out[54]= } \frac{4 \alpha}{\tau} + 8 M_0 (2 \beta - \gamma + 2 \gamma \tau) + \frac{M_0^2 (\delta + 2 \gamma \tau)}{\tau} + \frac{16}{3} (\gamma + 6 \beta \tau - 6 \gamma \tau + 8 \gamma \tau^2) + \frac{\alpha \text{Log}[N_x^2]}{\tau}$$

We see that the cost per processor actually scales with the number of processors, but only logarithmically due to the Allreduce. We do not expect to see this